INKLOOM AI / HACKATHON SUBMISSION

BrandProof

Round 2 | Build progress report

BrandProof is a human-led workspace for checking brand content against team rules and approved evidence, then recording review decisions against the exact content version. The Round 2 build turns the original proposal into a deployed full-stack application with text and document workflows, structured checks, review controls, and traceable history.

LIVE PROJECT	SOURCE CODE
https://brandproof-95p8.onrender.com	github.com/srushti234oss/BrandProofAI

TEAM	Muted Mavericks
TEAM MEMBERS	Sathwikaa, Soumya H, Srushti Shelar, Bhavana K
EVENT / PROBLEM	Inkloom AI - Intelligent Brand Governance
SUBMISSION STAGE	Round 2 - working build

LIVE DEMO NOTE	The hosted app is available, but its optional Groq analysis was returning HTTP 403 during the latest reported demo attempt. Deterministic rule checks remain available; a successful live AI call still needs to be confirmed after provider permissions are corrected.

01 What is completed

The core journey is implemented: a team sets up a workspace and brand guidance, submits content, reviews findings with their sources and explanations, makes a human decision, and can inspect the version and audit history.

FEATURE	ROUND 2 IMPLEMENTATION	STATUS
Accounts and workspaces	Registration/login, workspace membership and role-aware permissions.	**Implemented**
Brand rules and guidelines	Guidelines can be pasted or uploaded; structured rules can be managed and versioned.	**Implemented**
Approved evidence	Workspace evidence can be uploaded or entered, approved, retrieved and linked to claims.	**Implemented**
Content intake	Pasted text, TXT, PDF, DOCX, PNG and JPEG (up to 20 MB); extraction method, status and warnings are retained.	**Implemented**
Risk Radar	Deterministic terminology/absolute-claim checks; claim classifications; findings explain rules and evidence status. Missing evidence is labelled unverified, not false.	**Implemented**
Visual checks	Configured palette comparison, logo template matching and selected OCR/contrast estimates are advisory and depend on reference assets.	**Partial / configuration-dependent**
Human review and versioning	Comments, reviewer assignment, status history, approve/request changes/reject, correction drafts and re-review of new versions.	**Implemented**
Trust and governance	Printable Trust Receipt, claim ledger, review history, guideline clarity signals and team/invitation features.	**Implemented**
AI and embeddings	OpenAI-compatible structured analysis and embedding retrieval are implemented, but require a compatible provider/key. The live Groq call most recently returned HTTP 403.	**Partial / configuration-dependent**
Production persistence	Current Render demo uses temporary local SQLite/media; data can be lost on restart/redeploy.	**Not in current demo deployment**

Important honesty note

AI is an optional analysis layer, not the only review path. If the provider is unavailable or rejects a request, deterministic findings and unverified-claim triage remain; the warning does not mean AI analysis succeeded. Final publication decisions remain with a human reviewer.

02 Technology and implementation

The product is a single-origin web app. The browser UI and API are served from one deployed Django service, which keeps the demo setup simple.

LAYER	TECHNOLOGY / ROLE
Frontend	React, TypeScript, Vite and Tailwind CSS; single-page workspace interface.
Backend / API	Python, Django, Django REST Framework and Django ORM; session authentication, CSRF protection and role checks.
Data	SQLite for the current short demo; PostgreSQL supported through DATABASE_URL. Relational records cover workspaces, guideline versions, rules, assets, findings, decisions and audit events.
Document and image processing	PyPDF, python-docx, PyMuPDF, Tesseract OCR, Pillow and OpenCV for extraction and selected visual checks.
AI / retrieval	Optional OpenAI-compatible chat and embedding endpoints. Approved evidence and guideline chunks are selected for a review; vectors are stored in JSON and cosine-ranked in application code, with lexical fallback.
Deployment	Docker image builds the frontend and runs Django/Gunicorn plus an analysis worker on Render. The current free demo uses temporary filesystem storage.

Core data flow

1 SET UP	2 SUBMIT	3 EXTRACT	4 CHECK	5 REVIEW	6 RECORD
Rules, guidelines, approved evidence	Text or supported file	Text/OCR + metadata	Deterministic checks + optional AI + retrieval	Human decision and comments	Version hash, receipt and audit history

Current operational limits

- The hosted demo is temporary: accounts, reviews and uploaded files may disappear when the free service restarts or redeploys.
- Embedding retrieval works for small workspaces, but there is no pgvector database index yet.
- OCR, logo matching and contrast checks can be incomplete or wrong; they are advisory, not compliance certification.
- External AI/provider availability and permission setup must be verified in the live environment; latest observed provider response was HTTP 403.

03 What changed from the original plan

Round 1 described a planned product and candidate architecture. Round 2 focuses on a working, human-led governance workflow and makes the implementation boundaries explicit.

ORIGINAL PLAN	ROUND 2 REALITY	WHY / EFFECT
PostgreSQL as the primary database	PostgreSQL remains supported, but the short Render demo uses fresh SQLite on temporary disk.	Avoided the free-tier database constraint for a quick demo. Not suitable for persistent production use.
pgvector or Chroma for semantic search	Vectors are stored as JSON and cosine-ranked in application code; lexical retrieval remains available.	Kept retrieval functional at small scale; database-side indexing is deferred.
Multimodal review across text, images and documents	Text/PDF/DOCX and common image intake are implemented, with OCR and selected heuristic visual checks.	Visual analysis is deliberately advisory; no claim of complete multimodal understanding.
AI-assisted evidence-backed analysis	Structured LLM analysis is optional; rules remain active without it. The latest live test showed HTTP 403.	Provider permissions/key access need resolution before presenting a successful live AI call.
Production-ready retention and storage	Local/private storage and retention controls exist in the app; production S3, database backups and scheduled retention require host configuration.	The current hosted environment is a short-lived demo configuration.
Round 1 research and accuracy targets	No user research results or quantitative model-accuracy claims are included here.	Report only measurements after running a documented evaluation.

How to demonstrate the completed core path

- Register or sign in and create a workspace.
- Add one brand guideline and a structured rule; add an approved evidence passage.
- Submit a short piece of original text or a supported document.
- Open findings and show the matched rule/source, evidence status and cautious correction options.
- Make a human review decision, then show the saved version and Trust Receipt.

For a clean demo account, enter your own non-sensitive material. Do not rely on seeded sample records or imply a provider request succeeded while the HTTP 403 is unresolved.

04 Three meaningful Round 3 improvements

These are planned product and engineering additions, not cosmetic changes. Success measures below are targets for validation, not results already achieved.

1

Campaign-wide consistency review

Let a user link a campaign’s copy, image, PDF and supporting product facts. Compare claims, numbers, disclaimers and approved wording across the linked assets, then show conflicts with the exact source locations.

Measure: every reported conflict points to the two relevant passages/assets; evaluate on a labelled set of consistent and intentionally conflicting campaigns, tracking missed conflicts and false alarms.

2

Indexed hybrid evidence search

Move evidence and guideline embeddings into PostgreSQL with pgvector indexing, combine semantic and lexical retrieval, and preserve source/page provenance. Add retrieval evaluation so the team can check whether the right approved source was surfaced.

Measure: report retrieval precision/recall on a labelled query set and compare latency as the evidence collection grows; never return unapproved evidence.

3

Persistent, resilient team deployment

Replace temporary demo storage with managed PostgreSQL and private object storage; run the analysis worker separately with retries and clear job states; add backups, operational monitoring and a recovery runbook.

Measure: confirm records/files survive redeploys, queued work recovers after worker restart, and backup restoration is demonstrated in a staging environment.

Why these matter

Together, these improvements would let BrandProof compare a whole campaign, retrieve approved information more reliably as it grows, and keep team data and analysis jobs dependable between deployments. They extend the product’s core governance capability and reliability.

05 Submission and judging map

The submitted links and the current build status are collected here for reviewers.

WORKING PROJECT	https://brandproof-95p8.onrender.com
SOURCE REPOSITORY	https://github.com/srushti234oss/BrandProofAI
ROUND 2 PDF	This progress report

Judging criteria alignment

CRITERION	EVIDENCE IN THIS BUILD
Working core workflow - 30	Account/workspace setup, guideline and evidence configuration, content intake, findings, human decision and version-bound record.
Technical depth and reliability - 25	React/TypeScript + Django/DRF, ORM data model, structured AI integration, retrieval, OCR/visual processing, permission controls and documented deployment limits. The live Groq call currently needs its 403 permission issue resolved.
Progress from Round 1 - 20	The original proposal is now a deployed application with implemented workflows and explicit distinctions between implemented, advisory and configuration-dependent capabilities.
Three planned improvements - 15	Campaign-wide consistency, indexed hybrid retrieval, and persistent resilient deployment are described with validation measures.
Usability - 10	A reviewer can follow a clear setup-to-decision workflow. Live review should use a fresh workspace and non-sensitive content; Render free service may take time to wake.

Evaluation and claim boundaries

This document does not claim measured accuracy, external user-study results, production persistence, guaranteed OCR correctness, or a successful Groq analysis call. The current PDF and source repository describe the build as it stands; external service access remains dependent on configuration and account permissions.

BrandProof helps teams review content. It does not guarantee factual truth, legal compliance, accessibility compliance, or publication safety; an authorised person makes the final decision.