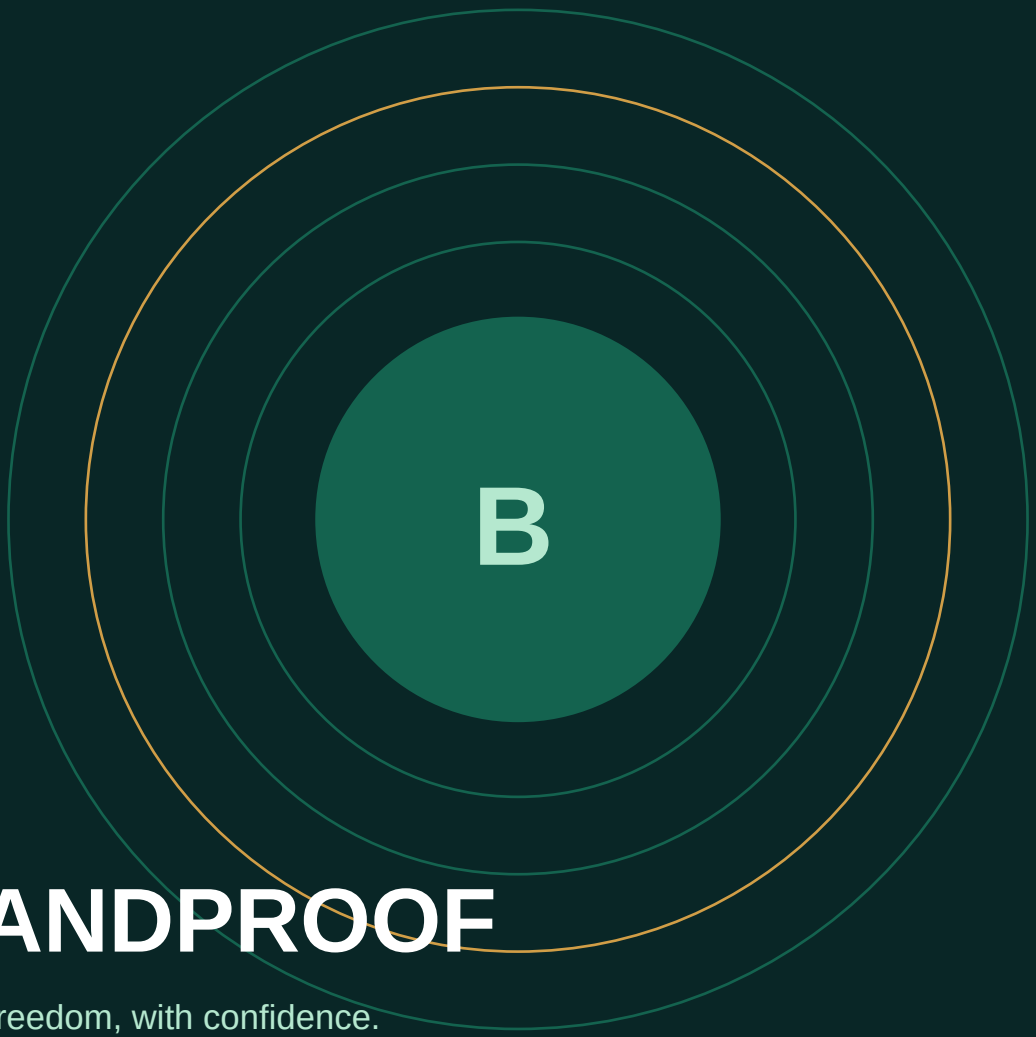

BRANDPROOF

Creative freedom, with confidence.

A human-led brand governance workspace that checks campaign content against the rules and evidence a team provides.

ROUND 3 / FINAL PRODUCT REPORT

MUTED MAVERICKS

Inkloom AI | 30 October 2026

LIVE DEMO

brandproof-95p8.onrender.com

SOURCE

github.com/srushti234oss/BrandProofAI

Submission report reflects implementation and owner-reported live checks as of 30 Oct 2026.

01 / THE PRODUCT

A review journey built for people

BrandProof helps teams review promotional content before publication. A reviewer can see what matched a rule, which claim needs evidence, and what action is still theirs to decide.

01

SET BRAND RULES

Add guideline text or files, structured rules, and approved evidence.

02

SUBMIT CONTENT

Create an asset, name its campaign, paste copy or upload a supported file.

03

CHECK & TRACE

Rules and optional Gemini analysis identify risks; evidence and source versions stay visible.

04

REVIEW & DECIDE

Assign, comment, request changes, approve or reject. A correction creates a new version.

WHAT IS INCLUDED

- Accounts, workspaces, roles, invitations and team collaboration
- Guideline library, versioning, structured terminology rules and approved evidence
- TXT, PDF, DOCX and image submission with OCR when available
- Deterministic checks plus optional Gemini-assisted analysis
- Claim ledger, reviewer assignment/comments, decisions, version history and Trust Receipt
- Campaign comparison, guideline clarity signals, and advisory palette/logo/contrast checks

A clear boundary

BrandProof organizes review signals; it does not certify factual truth, legal compliance, or publication safety. Authorized people make final decisions.

02 / IMPROVEMENTS

The three Round 2 commitments

These are the same three substantial improvements committed in Round 2. Progress is reported precisely: the first is demonstrated in the product, the second is configured with live PostgreSQL according to the owner, and the third is only partly complete.

01

Campaign-wide consistency review

DELIVERED: FIRST RELEASE

Related assets share a campaign name. The check compares readable latest versions for changed numbers in repeated sentence patterns, configured preferred wording, and recognized disclosure-marker presence. It displays the exact assets, versions and passages that differ.

User value: catch a mismatch such as 15% vs 20% before publication, then send it to a human reviewer.

Boundary: conservative pattern signals; not semantic paraphrase matching, truth verification, or legal/disclosure compliance.

02

Indexed hybrid evidence search

IMPLEMENTED; LIVE PATH OWNER-REPORTED

PostgreSQL is connected and the Settings screen reports hybrid pgvector HNSW plus lexical fusion. The code includes an EmbeddingIndex model, a pgvector HNSW cosine index migration, and reciprocal-rank fusion. Workspace, approval and guideline-version scope are rechecked; fallbacks remain for SQLite or unavailable vectors.

User value: retrieve approved, relevant brand evidence through semantic and keyword matching with provenance.

Verification still needed: independently exercise the deployed vector/embedding path and measure retrieval quality; current evidence is owner-reported UI plus code, not a recorded benchmark.

03

Persistent, resilient team deployment

PARTIAL: DATABASE CONNECTED; RESILIENCE NOT VERIFIED

The owner reports that PostgreSQL is connected and the server is live. A separate persistent Render profile is present in the repository, but the current deployment profile is the free single-service setup. Separate worker, private object storage, backups/restore, monitoring and end-to-end redeploy persistence remain to be deployed and verified.

User value when completed: team content and uploads survive restarts, while slow analysis runs independently and recovers from failures.

Honest status: do not describe this commitment as fully delivered until the persistent profile, uploads, worker and recovery checks pass on the public service.

03 / IMPLEMENTATION

How the system fits together

One workspace-scoped product flow connects a React client, Django APIs, relational records, document extraction and optional AI. Uploads and AI outputs are treated as review inputs, not final authority.

CLIENT	React, TypeScript and Vite; same-origin frontend bundled into the Django service.
API & ACCESS	Python, Django and Django REST Framework; session authentication, CSRF, workspace role checks and throttling.
DATA	SQLite for local development; PostgreSQL through DATABASE_URL for hosted database use.
AI	Gemini configured for analysis in the current deployment according to owner confirmation. Provider secrets belong in host environment settings.
RETRIEVAL	Guideline/evidence chunks with lexical ranking and semantic vectors; pgvector HNSW + reciprocal-rank fusion where enabled.
FILES	PDF/DOCX text extraction, image OCR with Tesseract where configured, and source provenance via Django storage.
ANALYSIS JOBS	Database-backed analysis state; an optional separate worker/persistent Render profile is described in render.persistent.yaml.

DATA FLOW

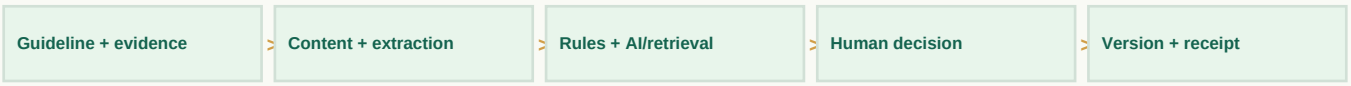

DEPLOYMENT LIMIT TO KEEP VISIBLE

The repository has two Render profiles: render.yaml is the temporary free single service; render.persistent.yaml defines the intended web + worker + PostgreSQL + private S3-compatible setup. The latter requires paid host resources and storage credentials. The report does not claim it is live.

04 / EVIDENCE

Reliability, checks & known gaps

Round 3 asks for a product that survives realistic use. This section distinguishes code-level checks from external hosting checks that still require a real deployed run.

RELIABILITY WORK IN THE REPOSITORY

- Dashboard panels load independently so one failed request does not blank the whole view.
- Malformed non-text content is rejected with a client error instead of crashing string handling.
- Duplicate review decisions are idempotent for the same reviewer and version.
- Repeated retry requests reuse the existing queued/processing analysis.
- Worker logic supports bounded retries with backoff and stale-job recovery.
- AI failure preserves deterministic findings and unverified-claim triage.

CHECKS RECORDED IN THE PROJECT NOTES

Frontend production build	PASS - TypeScript build and Vite bundle.
Django system + migrations	PASS - no system issues or pending model migration.
Focused campaign/retrieval/retry checks	PASS - campaign signals, workspace/approval scope, vector-index mirror and worker retries.
End-to-end API flow	PASS - registration through guideline/evidence, analysis, human review and Trust Receipt.
Backend suite	46 pass; one schema test blocked in the local environment by missing inflection dependency.

STILL TO VERIFY BEFORE CALLING THE THIRD IMPROVEMENT COMPLETE

- Deploy the persistent web + worker + PostgreSQL + private object-storage profile; confirm actual service health.
- Test file upload/download and content persistence through service restart and redeploy.
- Verify vector search with the deployed embedding provider; record a small relevance/latency evaluation.
- Configure database backups, perform a restore drill, and add health/worker/storage monitoring plus a runbook.
- Install backend requirements and rerun the one schema test blocked by missing local inflection.

05 / DELIVERY

Close-out plan & roadmap

The working app and source are available at the links below. The public deployment should be treated as a judge entry point, while the following checks close the infrastructure and verification gaps honestly identified above.

SUBMISSION LINKS

LIVE PROJECT

https://brandproof-95p8.onrender.com

SOURCE REPOSITORY

https://github.com/srushti234oss/BrandProofAI

NEXT ACTIONS

1. Confirm the final deployed commit matches the submitted repository.
2. Deploy and health-check the persistent profile; verify upload, restart and redeploy persistence.
3. Exercise the 15% vs 20% campaign conflict and a matching campaign; confirm precise asset/version context.
4. Run a retrieval relevance check against approved evidence; record result and provider/model settings.
5. Install requirements, run the complete backend suite, and publish clone/run instructions and recovery notes.

ROADMAP

NOW	Finish and verify the exact three commitments; close deployment and test gaps.
NEXT	Backup/restore drills, alerting, storage retention enforcement and an operator runbook.
LATER	Build a labelled campaign benchmark for paraphrase and disclosure variants; measure retrieval precision/recall; expand reviewer export and accessibility checks.

A promise grounded in evidence

BrandProof keeps the evidence, rule, version and human decision connected. It helps a team review with confidence; it never replaces the reviewer.