INKLOOM AI / ROUND 3 SUBMISSION

BrandProof

A human-led brand governance workspace: compare content with team rules and approved evidence, review findings, and record decisions against the exact content version. This report says what is implemented, what passed local checks, and what still depends on provider or hosting setup.

Report date: 29 October 2026
The code work described below is in the current working copy. The public Render service has not yet been updated from this copy. The latest user-shared production screenshot showed Groq HTTP 403. The optional paid persistent profile is prepared but not deployed.

SUBMISSION ITEM	DETAIL
Live project	brandproof-95p8.onrender.com (current demo URL)
Source repository	github.com/srushti234oss/BrandProofAI
Team	Muted Mavericks - Sathwikaa, Soumya H, Srushti Shelar, Bhavana K
Event	Inkloom AI - Intelligent Brand Governance

Judge path

Register and create a workspace; add a real guideline, structured rule and approved evidence; submit text or a supported file; inspect findings; record a human decision; open the versioned Trust Receipt. The app does not seed sample workspace data. Use a fresh account rather than treating the earlier BrightBite screenshot as verified evidence.

Status terms

Implemented means code is present. **Locally verified** means checks passed in this environment. **Pending live setup** means judge-facing deployment/provider behavior still needs owner configuration and a fresh check.

01 / Features implemented

Core account-to-review flow

FEATURE	WHAT WORKS	LIMIT / STATUS
Accounts/workspaces	Registration, login, workspace membership, invitations, role-aware permissions.	SMTP optional; manual invitation links work.
Brand guidelines	Paste/upload guideline source, version history, create/edit rules, approved examples.	Document extraction may return warnings.
Evidence and claims	Add/upload evidence, approve/revoke it, retrieve approved workspace evidence, link sources to claims, reviewer adjudication.	Missing evidence stays unverified; it is not called false.
Content intake	Pasted copy, TXT/PDF/DOCX/PNG/JPEG to 20 MB; extraction metadata, warnings, source retention.	OCR needs Tesseract.
Risk Radar	Deterministic term/claim rules, optional structured LLM analysis, cautious correction suggestions.	AI is advisory; latest user screenshot showed Groq HTTP 403.
Campaign Consistency	Across latest readable versions in a named campaign: changed numbers in repeated sentence patterns, configured preferred-wording variants, recognized disclaimer-marker presence. Exact asset/version/wording shown.	Conservative signals for human review. Does not judge truth, legal compliance, paraphrases, or format-specific disclosure requirements.
Review and audit	Comments, reviewer assignment, status history, approve/request changes/reject, corrections as new versions, Trust Receipt and audit history.	Approval stays bound to its exact content version.
Clarity and visual checks	Repeated review signals can suggest guideline clarification; palette/logo/selected contrast checks run when references exist.	Simple signal threshold; visual results are heuristic/advisory.

Core data flow

Configure guidelines and approved evidence → submit content → extract text and preserve provenance → deterministic checks plus optional AI/retrieval → human resolves findings and decides → version hash, receipt and audit history.

No demo seeding

The standard Render startup only runs database migrations. The team or judge creates an account and enters its own workspace information. `seed_demo` is an optional local walkthrough command and is not run by either Blueprint.

02 / Three Round 2 improvements

Progress against the exact commitments

1. Campaign-wide consistency review

WORKING FIRST RELEASE IMPLEMENTED

A new workspace-scoped screen/API compares latest readable versions of same-name campaigns. It flags numeric differences in repeated sentence structures, presence/absence of configured preferred wording, and presence/absence of recognized disclaimer markers. It shows both passages and versions. It remains a rules-based, conservative check: no semantic paraphrase comparison, legal decision, or determination of which claim is correct.

2. Indexed hybrid evidence search

IMPLEMENTED IN CODE; LIVE PGVECTOR PATH NOT VERIFIED

Added an EmbeddingIndex model, a PostgreSQL migration that creates the pgvector HNSW cosine index, and reciprocal-rank fusion of semantic and lexical retrieval. Retrieval rechecks current workspace, approval and guideline version. SQLite/unavailable vector cases retain JSON cosine and lexical fallbacks. Database-side ranking needs PostgreSQL/vector extension and a compatible embeddings provider returning 1536 dimensions.

3. Persistent, resilient deployment

SEPARATE PAID PROFILE PREPARED; NOT LIVE

`render.persistent.yaml` defines separate web and analysis-worker services with a dedicated PostgreSQL database and private S3-compatible upload storage. It applies migrations at web startup, shares DB/storage config with the worker, and enables up to three analysis attempts with backoff plus stale-job recovery. It requires paid Render resources and external storage credentials; it has not been created or tested. Backups, restore and monitoring still need owner setup and verification.

External steps still required: push this working copy to GitHub, deploy it, configure the isolated database/storage/provider, then run the judge journey on the new URL. The existing live demo link may still serve the old commit and temporary data.

03 / Technical architecture

Single-origin React + Django application. The Docker build bundles the frontend into the Django service; session authentication and CSRF keep browser/API calls on one origin.

LAYER	IMPLEMENTATION
Frontend	React, TypeScript, Vite; client-side workspace and review screens.
API	Python, Django, Django REST Framework; ORM, session auth, CSRF, role checks, request throttling.
Database	SQLite by default locally and in temporary demo; PostgreSQL through DATABASE_URL in durable profile.
AI	Optional OpenAI-compatible chat and embedding endpoints; secrets supplied by host environment.
Retrieval	Approved evidence and current guideline chunks. Hybrid semantic + lexical reciprocal-rank fusion. pgvector HNSW for supported PostgreSQL vectors; portable JSON/lexical fallback.
Documents/images	PyPDF, python-docx, PyMuPDF, Tesseract, Pillow, OpenCV; original source through Django storage.
Jobs	Database-backed AnalysisRun state machine; retries/backoff and stale processing recovery in worker command.
Render profiles	render.yaml: free single service and ephemeral SQLite/media. render.persistent.yaml: paid web + separate worker + PostgreSQL + private S3-compatible media.

Workspace protections

- Evidence retrieval filters to current workspace and approved source records; pgvector candidates are rechecked before use.
- Guideline retrieval filters to current guideline version.
- Corrections make a new version; old reviewer decisions remain on the old version.
- AI/provider failure retains deterministic findings and unverified claim triage instead of implying compliance.

04 / Reliability and testing

Hardening added

- Dashboard loads panels independently; one failed API request no longer blanks all dashboard content.
- Malformed non-text asset fields return HTTP 400 rather than causing string-operation errors.
- Repeated identical review decisions are idempotent for the same reviewer and version.
- Repeated retry click returns the existing queued/processing job.
- Worker can retry a failed analysis up to three times with backoff; stale processing work is reclaimed.

Checks run

CHECK	RESULT
Frontend production build	PASS - TypeScript build and Vite bundle.
Django system check	PASS - no issues.
Migration consistency	PASS - no model migration pending.
Focused campaign/retrieval/retry checks	PASS - campaign numeric/wording/disclosure, workspace approval scope, vector index mirror, and worker retries.
End-to-end API flow	PASS - fresh registration, workspace, guideline/rule, approved evidence, content analysis, review and Trust Receipt.
Runnable backend suite	46 tests pass; the one remaining schema test cannot run here because global Python lacks `inflection`. It is listed in backend requirements. Install `pip install -r requirements.txt` to run all tests.

What is not verified yet

- Full local suite: one API schema test blocked by missing local `inflection`; not a product code failure, but still needs rerun after dependency installation.
- Hosted service: latest user screenshot shows Groq HTTP 403. Enable key/model access and run a fresh review.
- PostgreSQL/pgvector and embedding provider integration have not been exercised against the hosted database.
- Persistent profile has not been deployed; S3 credentials, backup/restore and health monitoring are not configured.
- Free demo persists neither database nor uploads across restarts/redeploys.

05 / Changes, roadmap and submission checklist

Changes from the original proposal

Round 1 proposed PostgreSQL/pgvector, broader campaign comparison, multimodal analysis and production-ready retention. Round 2 delivered the human-led workflow and disclosed that free demo storage, provider permission and pgvector hosting were incomplete. Round 3 adds a first campaign consistency product surface, indexed hybrid search code, and an opt-in durable infrastructure profile. The live external setup remains a required team action; no production success is claimed before it is tested.

Before submitting

- Commit and push this working copy; confirm the final GitHub repository link points to that commit.
- Choose whether to create the paid Render profile; check current service/database and storage costs first.
- Add the isolated S3-compatible bucket and Groq secrets in Render, then wait for web, worker and database health.
- Run a fresh account journey; compare two campaign assets; upload/download a private file; review and open a receipt; refresh and verify persistence.
- After fixing Groq 403, create new content and confirm the analysis warning clears; verify AI and pgvector settings explicitly.
- Install backend requirements and rerun all tests; inspect Render build/deploy logs.
- Update the final URL here and in the submission form only after successful deployment.

Roadmap

STAGE	PRODUCT PLAN
Round 3 close	Validate numeric, preferred-wording and disclaimer comparison on real campaign examples; validate hosted hybrid retrieval; deploy persistent profile and test worker restart/redeploy/file persistence.
Next	Configure and test DB backups/restoration, S3 lifecycle/retention, worker/web/database alerts, and document an incident/recovery runbook.
Later	Build a labelled campaign benchmark for paraphrases and format-specific disclosures; measure retrieval precision/recall and latency; expand onboarding and reviewer exports.

BrandProof supports human review. It does not certify factual truth, legal compliance, accessibility compliance or publication safety. Final decisions belong to an authorised reviewer.

Project: current demo URL | Repository: BrandProofAI